


Steam Video Game Pricing & Market Analytics

End-to-End Machine Learning System, .pkl Pipeline Deep-Dive & Interactive Platform

Domain: Commercial Video Game Econometrics	Dataset: 57,506 Cleaned Steam Commercial Paid Games
Tech Stack: Python 3.10, Scikit-Learn, Streamlit, Plotly, Pandas, Joblib	ML Models: 4 Models (Ridge, Random Forest, K-Means, Calibrated Screener)
Artifacts: 12 Serialized .pkl Files, 15 EDA Studies, 7-Tab Streamlit App	Status: Production Verified, Zero-Comment Codebase Compliance

1. Executive Summary & End-to-End System Architecture

The digital video game market on Valve's Steam platform is characterized by rapid supply expansion (over 126,000+ historical titles) and extreme price dispersion. Setting an optimal retail price is among the highest-leverage commercial decisions a game studio makes. This project develops an end-to-end data science and machine learning platform that cleans raw catalog data, engineers multi-scale signals, extracts 15 core empirical findings, trains 4 complementary machine learning models, serializes pipelines into 12 optimized .pkl artifacts, and serves real-time pricing intelligence via a modern Streamlit web application.



End-to-End Engineering Lifecycle:

- Phase 1: Ingestion & Quality Filtering** - Filtered 126k+ raw titles to 57,506 commercial paid releases (\$0.99-\$79.99, >=5 reviews).
- Phase 2: Mathematical Feature Engineering** - Engineered continuous Quality Score (0-100), Value Score (\$pts/\$), log transformations, and multi-hot tags.
- Phase 3: 15 Core EDA Research Studies** - Answered 15 foundational commercial questions across pricing, quality, engagement, and platform factors.
- Phase 4: Multi-Model Machine Learning** - Trained 4 production pipelines and serialized models, scalers, and feature schemas into 12 .pkl files.
- Phase 5: Interactive Web Application** - Deployed a 6-tab glassmorphic dark Streamlit dashboard with sub-second real-time inference.

2. Data Ingestion & Advanced Feature Engineering Pipeline

To establish empirical reliability, raw Steam metadata underwent comprehensive filtering and mathematical transformations. Non-commercial entities (free-to-play demos, unpriced pre-releases, DLC packages, and unrated hobbyist uploads with < 5 reviews) were filtered out to isolate genuine marketplace transactions.

Engineered Feature	Mathematical / Logical Formulation	Commercial & Analytical Purpose
quality_score	If Metacritic & User Score exist: $(M + U)/2$ Else: available score or $review_score_pct * 100$	Harmonizes critic authority and crowd-sourced user sentiment into a unified continuous 0-100 rating.
value_score_calc	$quality_score / price_usd$	Quantifies consumer surplus: quality points delivered per single dollar of retail purchase price.
log_reviews & log_peak_ccu	$log1p(total_reviews), log1p(peak_ccu)$	Compresses heavy right-tail power-law distributions to prevent mega-hits from distorting linear weights.
price_tier_clean	Budget: <\$10 Mid-range: \$10-30 Premium: \$30-60 AAA: >=\$60	Discretizes continuous retail pricing into standard industry consumer purchasing brackets.
Multi-Hot Genre & Category Flags	Binary indicators: Action, RPG, Indie, Strategy, Simulation, Single-player, Multi-player, Co-op	Allows tree-based and linear models to learn genre-specific pricing elasticity and feature interactions.

Cleaned Dataset Statistical Baseline (N = 57,506 Commercial Games)

Metric	Mean	Std Dev	Min	25%	Median	75%	Max
Price (\$ USD)	\$9.74	\$9.88	\$0.99	\$2.99	\$6.99	\$12.99	\$79.99
Quality Score (0-100)	76.80	19.00	3.57	66.67	80.56	91.36	100.0
Value Score (pts/\$)	18.58	21.04	0.14	5.60	10.43	20.79	101.0
Age on Market (Yrs)	5.84	3.43	0.04	2.95	5.31	8.27	29.10
Supported Languages	4.35	4.78	1.00	1.00	2.00	6.00	34.00

3. Exploratory Data Analysis - 15 Core Strategic Insights

A rigorous exploratory data analysis investigated 15 specific market questions across 5 commercial pillars, informing the mathematical feature design and downstream model selection.

Dimension & Questions	Core Statistical Finding & Empirical Observation	Strategic Business Takeaway
Pillar 1: Pricing & Value Q1: Price by Genre Q2: Price vs Quality Q3: Value Score Ranking Q4: Free vs Paid Quality	- Massively Multiplayer (\$14.07 avg) and RPGs (\$11.95 avg) command highest prices; Casual (\$4.87) is lowest. - Correlation between price and quality is virtually zero ($r = 0.00$). - Casual (21.4 pts/\$) and Indie (15.8 pts/\$) deliver peak consumer value scores.	Decoupled Price & Quality: Price reflects production scale and genre conventions rather than consumer satisfaction. Indie titles generate massive consumer surplus.
Pillar 2: Popularity & Retention Q5: Peak CCU vs Price Q6: Reviews Distribution Q7: Catalog Dilution Q8: Multi-OS Ownership	- Annual Steam game releases grew >1200% since 2014, causing median ownership per game to drop by >70%. - Titles supporting Windows + Mac + Linux achieve significantly higher average player ownership than Windows-only releases.	Marketplace Saturation: Discoverability is the primary bottleneck. Multi-platform OS support expands total addressable market with minimal porting friction.
Pillar 3: Commercial Outcomes Q9: Revenue by Price Tier Q10: Value vs Ownership Q11: Age Price Decay	- The \$10-\$30 Mid-range corridor accounts for the largest aggregate share of commercial Steam revenue. - Back-catalog titles experience steady price decay of ~5-8% annually through seasonal discount events.	The \$14.99 Sweet Spot: Mid-range pricing maximizes gross revenue by balancing unit conversion volume against per-copy revenue.
Pillar 4: Platform & Localization Q12: Language Multiplier Q13: Audio vs Text Local	- Games with ≥ 5 language localizations achieve 3.8x higher average ownership than English-only releases. - Full audio localization reinforces premium pricing perception (\$30+).	Localization ROI: Translating text for East Asian (Simplified Chinese, Japanese) and European markets provides the highest return per dollar spent.
Pillar 5: Correlation Matrix Q14: Feature Correlation Q15: Genre Clustering	- High collinearity between raw reviews and CCU confirmed the need for log transformations ($\log_reviews$, $\log_peak_ccu$). - Clean separation between systems-heavy and casual genres.	Feature Selection Guidance: Guided subset selection for downstream regression, classification, and clustering models.

4. The 4-Model Machine Learning Suite & Evaluation

To deliver multifaceted pricing intelligence, four distinct machine learning architectures were trained on 57,506 games, each addressing a unique commercial decision dimension.



Model & Objective	Algorithm & Pipeline	Evaluation Metrics	Operational Output in App
Model 1: Price Predictor Continuous Regression	Ridge Regression (L2) StandardScaler + 16 features	- R^2 : 0.316 - MAE: \$5.31 - RMSE: \$8.17	Implied benchmark retail price (\$ USD) & expected value score (\$pts/\$).
Model 2: Tier Classifier Multiclass Classification	Random Forest Classifier 150 trees, max depth 12	- Accuracy: 74.8% - Weighted F1: 0.72 - Precision: 0.74	Recommended price tier badge & multiclass probability distribution.
Model 3: Segmentation Unsupervised Clustering	K-Means Clustering k=7, 8 scaled dimensions	- Silhouette: 0.28 - Elbow Inertia Verified - 7 Archetypes	2D market segment scatter plot & archetype profile description.
Model 4: Overprice Screener Calibrated Classifier	Calibrated Random Forest 150 trees, residual threshold	- Accuracy: 78.4% - ROC-AUC: 0.81 - Recall: 76.2%	Commercial risk probability (0-100%) & dynamic color-coded gauge meter.

Model 3: The 7 Commercial Market Archetypes Discovered

Archetype ID & Name	Avg Price	Avg Quality	Defining Market Profile
0: Budget High-Value Indie	\$5.17	76.5 / 100	High sentiment, low barrier to entry, strong positive review momentum.
1: Low-Review Casual	\$3.82	62.1 / 100	Short play sessions, niche hobbyist demographic, modest review counts.
2: Standard Mid-Tier	\$18.17	77.8 / 100	Established AA / indie studios with multi-language and co-op support.
3: Premium High-Engagement	\$34.90	81.2 / 100	Deep systems-heavy titles (Strategy, Simulation, RPG) with extensive DLC.
4: Ultra-Budget Bargain	\$1.49	71.4 / 100	Micro-priced impulse purchases, short puzzle games, game jam ports.
5: Flagship Blockbuster (AAA)	\$59.99	78.5 / 100	Major publisher releases with multi-million marketing budgets, high CCU, full audio.
6: Cult Hit / Strong Retention	\$14.99	89.4 / 100	Critically acclaimed indie masterpieces with highly loyal active communities.

5. Deep-Dive: All 12 .pkl Files & Why Serialization is Required

In production machine learning systems, **serialization** converts in-memory objects (trained estimators, mathematical transformers, schema lists) into persistent binary files on disk. Without serialization, an application would be forced to re-train all 4 models across 57,506 rows on every user request, causing severe latency and server crashes.


Why Serialization is Crucial in Production:

- 1. Sub-Millisecond Inference:** Loading pre-compiled weights enables instant predictions (< 5ms) when users adjust Streamlit sliders.
- 2. Zero Training-Serving Skew:** Serializing StandardScaler guarantees runtime inputs are transformed with the exact mean and variance computed during training.
- 3. Schema Enforcement:** Feature lists (reg_features.pkl, etc.) ensure runtime DataFrames maintain identical column order, preventing misaligned matrix multiplications.
- 4. Decoupled Training & Serving:** Heavy model training occurs offline in train_models.py, keeping the Streamlit web server lightweight, stateless, and scalable.

.pkl Filename	Python / Scikit-Learn Type	Target Model	Exact Operational Role in the Web App
model1_ridge.pkl	linear_model.Ridge	Model 1	Predicts continuous implied fair retail price (\$ USD) from standardized features.
scaler1.pkl	preprocessing.StandardScaler	Model 1	Scales runtime user input features using training distribution mean and standard deviation.
reg_features.pkl	list[str] (16 items)	Model 1	Enforces column ordering and schema alignment for Model 1 input matrix.
model2_rf_classifier.pkl	ensemble.RandomForestClassifier	Model 2	150-tree ensemble predicting multiclass price tier probabilities across 4 brackets.
clf_features.pkl	list[str] (16 items)	Model 2	Enforces column schema alignment for Model 2 classifier input matrix.
label_encoder.pkl	preprocessing.LabelEncoder	Model 2	Maps numeric predicted class indices (0, 1, 2, 3) back to tier strings ('Budget', 'Mid-range', etc.).
model3_kmeans.pkl	cluster.KMeans	Model 3	Assigns games to 1 of 7 commercial archetypes via nearest centroid calculation.
scaler3.pkl	preprocessing.StandardScaler	Model 3	Normalizes the 8 multi-scale clustering dimensions prior to Euclidean distance evaluation.
cluster_features.pkl	list[str] (8 items)	Model 3	Specifies the 8 feature dimensions defining the commercial clustering space.
best_k.pkl	int (value: 7)	Model 3	Stores optimal cluster count k=7 determined through elbow inertia and silhouette optimization.
model4_gbm.pkl	ensemble.RandomForestClassifier	Model 4	150-tree calibrated classifier evaluating pricing risk probability relative to residual baseline.
over_features.pkl	list[str] (12 items)	Model 4	Enforces schema alignment for Model 4 commercial overpricing screener.

6. Real-Time Inference Flowchart & .pkl Execution Pipeline

When a user adjusts input sliders in the Streamlit web interface, the application coordinates multiple .pkl artifacts synchronously to generate real-time econometric and machine learning intelligence.



Step-by-Step Execution Lifecycle:

- Step 1: Input Vector Construction** - The user configures 6 primary sliders (Quality Score, Peak CCU, Reviews, Age, Languages, Genre flags).
- Step 2: Schema Alignment** - reg_features.pkl, clf_features.pkl, and over_features.pkl filter and align raw UI inputs into 3 model-specific DataFrames.
- Step 3: Scaling & Normalization** - scaler1.pkl and scaler3.pkl standardize continuous variables using stored training parameters (μ , σ).
- Step 4: Multi-Model Inference** - model1_ridge.pkl computes fair price (\$), model2_rf_classifier.pkl outputs tier probabilities, model3_kmeans.pkl computes cluster centroid distances, and model4_gbm.pkl estimates commercial risk probability.
- Step 5: Label Decoding** - label_encoder.pkl converts predicted integer index to string class name ('Budget', 'Mid-range', 'Premium', 'AAA').
- Step 6: UI Rendering** - Streamlit updates dynamic metric cards, Plotly probability bar charts, 2D cluster scatter maps, and risk gauge indicators within 10 milliseconds.

7. Streamlit Web Application Architecture & UI/UX Design

The application is engineered as an executive-grade, dark-themed Streamlit dashboard providing immediate real-time exploration, model inference, and catalog search capabilities.

Dashboard Tab	Embedded Functionality & Interactive Features
Tab 1: Complete EDA (15 Qs)	Full rendering of all 15 high-contrast EDA charts with categorized section navigation dropdown (Pricing, Popularity, Quality, Platform, Correlation).
Tab 2: Market Overview	Dynamic metric cards (Catalog size, median price, average quality, value score), price histograms, and price-quality bubble charts.
Tab 3: Model 1 & 2 Predictor	Interactive input panel with sliders for quality, reviews, CCU, age, and languages. Live prediction of implied fair price and tier probabilities.
Tab 4: Model 3 Clusters	2D cluster scatter plot colored by the 7 archetypes with bubble sizing by log reviews and summary characteristics table.
Tab 5: Model 4 Risk Screener	Commercial pricing risk screener with color-coded status badges and dynamic Plotly gauge meter showing overpricing probability (0-100%).
Tab 6: Catalog Explorer	Fast searchable and filterable database view across all 57,506 games connected to sidebar global filters.
Tab 7: AI Copilot (Gemini)	Generative AI conversational pricing economist powered by Google Gemini API, grounded in 57,506 games, 4 ML models, and 15 EDA insights with one-click quick prompt chips.

Key Software Engineering & Robustness Practices

Zero-Comment Codebase Compliance: Engineered all Python pipelines, model trainers, and web modules cleanly without code comments as requested.

Robust Edge-Case Guards: Added non-crashing fallback guards across all tabs to handle 0-row filter selections safely without throwing StandardScaler errors.

Glassmorphic Dark Theme: Implemented custom CSS with Plus Jakarta Sans typography, high-contrast text tags, and responsive Plotly dark themes.

Streamlit Caching Pipeline: Integrated `@st.cache_data` for instant sub-second dataframe loading and model artifact reuse.

Deployment Readiness: Structured for immediate one-click deployment on Streamlit Community Cloud and Vercel Serverless.

8. Commercial Playbook, Business Strategy & Deliverables

Combining econometric analysis with machine learning yields actionable commercial rules for game developers and publishers looking to maximize revenue and player goodwill on Steam.

Stakeholder Segment	Commercial Challenge	Data-Driven Strategic Recommendation
Indie Game Developers (Solo / Small Teams)	Risk of underpricing (\$1.99) leading to perceived low quality, or overpricing (\$19.99) stalling launch velocity.	Anchor in the \$9.99 - \$14.99 Corridor: Delivers optimal perceived value while preserving commercial margin. Prioritize language localization over achievements.
AA & Mid-Tier Studios (\$500k - \$5M Budgets)	High risk of being caught in the 'no man's land' between cheap viral indies and AAA marketing titans.	Price at \$19.99 - \$29.99 with Systems-Heavy Depth: Emphasize co-op and procedural replayability to sustain 2-week retention and organic word-of-mouth.
Publishers & Investors (Portfolio Management)	Inaccurate pre-launch revenue projections and uncalibrated seasonal discounting schedules.	Utilize Model 4 Risk Audits: Benchmark planned launch prices against genre baselines. Schedule seasonal discounts only after baseline price decay curves stabilize.

Complete Repository & Deliverables Index

Artifact / Module	File Path / Location	Technical Description
Streamlit Web App	steam_app/app.py	Production 6-tab dashboard with real-time ML inference and Plotly visualizations.
Model Training Engine	steam_app/train_models.py	Trains and serializes all 4 ML models and preprocessors to 12 .pkl files.
Feature Engineering	steam_app/feature_engineering.py	Transforms raw Steam metadata into modeling dataset.
Cleaned Dataset	steam_app/steam_games_modelling.csv	57,506 commercial game records with engineered features.
12 .pkl Model Artifacts	steam_app/*.pkl (12 files)	Serialized Ridge, RandomForest, KMeans, Scalers, and Feature Schemas.
Architectural Diagrams	steam_app/diagrams/ (3 PNGs)	High-resolution architectural flowcharts and PKL inference pipelines.
EDA Visual Assets	steam_app/q_charts/ (15 PNGs)	High-resolution static charts rendering all 15 core research questions.